



---

# Kubernetes: Container Orchestration and Micro-Services

University of Washington 590s  
2016-11-16

Alexander Mohr <[mohr@google.com](mailto:mohr@google.com)>  
Technical Lead / Manager on Google Container Engine and Kubernetes  
Github: [@alex-mohr](https://github.com/alex-mohr) Email: [mohr@google.com](mailto:mohr@google.com)

# Contents

1. Systems Projects at Google Seattle and Kirkland (2-3 mins)
2. Brief Docker Container Primer (5-10 mins)
3. Kubernetes: Container Orchestration (many mins)



# Prelude: Systems Projects at Google Seattle and Kirkland

## Seattle:

- Chrome Cloud (incl. Flywheel)
  - (Matt Welch)
- Flume / Dataflow / Apache Beam
  - (Craig Chambers)
- Compute Engine VM Hypervisor
  - (Mike Dahlin)
- Kubernetes + Container Engine
  - (Alex Mohr)
- App Engine Flex
  - (Tomas Isdal)
- Cloud Storage
  - (?)
- \$FOO
  - (Michael Piatek)

## Kirkland:

- Cloud Machine Learning
  - (Mona Attariyan)
- Spanner
  - (?)
- Compute Engine's Control Plane
  - (Mike Dahlin)
- Compute Engine's Persistent Disk
  - (?)
- Thialfi notifications
  - (Atul Adya)

These are some of the (public) projects explicitly focused on systems. Other areas require systems knowledge too!

# Contents

1. Prelude: Systems Projects at Google Seattle and Kirkland
2. **Brief Docker Container Primer**
  - a. Runtime
  - b. Building Images
  - c. Shipping Images
3. Kubernetes: Container Orchestration

# What are Containers? (Part 1: the Runtime)

Virtualize the kernel's syscall interface


- no guest OS or hypervisor as with VMs

Isolation (from each other and from the host)

- chroots
- namespaces
- cgroups

Packaging

- hermetically sealed bundles
- no external dependencies
- no DLL hell
- portable from dev laptop to on-prem & clouds



# What are Containers? (Part 2: Building an Image)

```
% cat - > Dockerfile
```

```
FROM node:4.4
```

```
EXPOSE 8080
```

```
COPY server.js .
```

```
CMD node server.js
```



# What are Containers? (Part 2: Building an Image)

```
% cat Dockerfile
FROM node:4.4
EXPOSE 8080
COPY server.js .
CMD node server.js
% docker build -t gcr.io/mohr-dev/hello-node:v1 .
[log spam]
```



# What are Containers? (Part 2: Building an Image)

```
% cat Dockerfile
FROM node:4.4
EXPOSE 8080
COPY server.js .
CMD node server.js
% docker build -t gcr.io/mohr-dev/hello-node:v1 .
[log spam]
% docker run -d -p 8080:8080 --name hello_tutorial gcr.io/mohr-dev/hello-node:v1
```



# What are Containers? (Part 2: Building an Image)

```
% cat Dockerfile
FROM node:4.4
EXPOSE 8080
COPY server.js .
CMD node server.js
% docker build -t gcr.io/mohr-dev/hello-node:v1 .
[log spam]
% docker run -d -p 8080:8080 --name hello_tutorial gcr.io/mohr-dev/hello-node:v1
% curl http://localhost:8080/
Hello World!
```



# What are Containers? (Part 3: Shipping an Image)

## The magic:

```
% gcloud docker --authorize-only  
% docker push gcr.io/mohr-dev/hellonode:v1  
The push refers to a repository [gcr.io/mohr-dev/hellonode] (len: 1)  
[...]  
v1: digest: sha256:d2f8b1387c535de6d6752a7c02c107576e86f9435d275be861fa8c6df5a29c4d size: 12985
```



# What are Containers? (Part 3: Shipping an Image)

## The magic:

```
% gcloud docker --authorize-only
```

```
% docker push gcr.io/mohr-dev/hellonode:v1
```

The push refers to a repository [gcr.io/mohr-dev/hellonode] (len: 1)

[...]

```
v1: digest: sha256:d2f8b1387c535de6d6752a7c02c107576e86f9435d275be861fa8c6df5a29c4d size: 12985
```

## Then, from any other machine:

```
% docker pull gcr.io/mohr-dev/hellonode:v1
```

v1: Pulling from mohr-dev/hellonode

```
Digest: sha256:d2f8b1387c535de6d6752a7c02c107576e86f9435d275be861fa8c6df5a29c4d
```

Status: Image is up to date for gcr.io/mohr-dev/hellonode:v1

```
% docker run $ARGS gcr.io/mohr-dev/hellonode:v1
```

...

# Contents

1. Prelude: Systems Projects at Google Seattle and Kirkland
2. Brief Docker Container Primer
- 3. Kubernetes: Container Orchestration**





Image by Connie Zhou

# Failures



A 2000-machine cluster will have  
1 to 10 machine failures per day.  
**This is not a problem: it's normal.**

# Kubernetes


Greek for “*Helmsman*”; also the root of the words “*governor*” and “*cybernetic*”

- Manages container clusters
- Inspired and informed by Google’s experiences and internal systems
- Supports multiple cloud and bare-metal environments
- Supports multiple container runtimes
- **100% Open source**, written in Go


Manage applications, not machines



# All you really care about



# The 10000 foot view



# Container clusters: A story in two parts



# Container clusters: A story in two parts

## 1. Setting up the cluster

- Choose a cloud: GCE, AWS, Azure, Rackspace, on-premises, ...
- Choose a node OS: CoreOS, Atomic, RHEL, Debian, CentOS, Ubuntu, ...
- Provision machines: Boot VMs, install and run kube components, ...
- Configure networking: IP ranges for Pods, Services, SDN, ...
- Start cluster services: DNS, logging, monitoring, ...
- Manage nodes: kernel upgrades, OS updates, hardware failures...

**Not** the easy or fun part, but unavoidable

This is where things like **Google Container Engine (GKE)** really help



# Container clusters: A story in two parts

## 2. Using the cluster

- Run Pods & Containers
- ReplicaSets & Deployments & DaemonSets & StatefulSets
- Services & Volumes & Secrets & Autoscalers

This is the fun part!

A distinct set of problems from cluster setup and management

Don't make developers deal with cluster administration!


Accelerate development by focusing on the applications, not the cluster



# Kubernetes: a Cloud OS?

Perhaps grandiose, but attempts at “Cloud OS” primitives:

- **Scheduling:** Decide where my containers should run
- **Lifecycle and health:** Keep my containers running despite failures
- **Scaling:** Make sets of containers bigger or smaller
- **Naming and discovery:** Find where my containers are now
- **Load balancing:** Distribute traffic across a set of containers
- **Storage volumes:** Provide data to containers
- **Logging and monitoring:** Track what's happening with my containers
- **Debugging and introspection:** Enter or attach to containers
- **Identity and authorization:** Control who can do things to my containers



# Workload Portability




# Workload portability

**Goal: Avoid vendor lock-in**

Runs in many environments, including  
“bare metal” and “your laptop”

The API and the implementation are  
100% open

The whole system is modular and  
replaceable



# Workload portability


**Goal: Write once, run anywhere\***

Don't force apps to know about concepts that are cloud-provider-specific

Examples of this:

- Network model
- Ingress
- Service load-balancers
- PersistentVolumes

\* approximately




# Workload portability

## Result: Portability

Build your apps on-prem, lift-and-shift  
into cloud when you are ready

Don't get stuck with a platform that  
doesn't work for you


Put your app on wheels and move it  
whenever and wherever you need




# Networking




# Docker networking




# Docker networking



# Port mapping



# Port mapping



# Kubernetes networking

## IPs are **cluster-scoped**

- vs docker default private IP

## Pods can reach each other directly


- even across nodes

## No brokering of port numbers


- too complex, why bother?

## This is a fundamental requirement

- can be L3 routed
- can be underlayed (cloud)
- can be overlayed (SDN)



# Kubernetes networking



# Pods



# Pods

**Small group of containers & volumes**

**Tightly coupled**

The atom of scheduling & placement


Shared namespace

- share IP address & localhost
- share IPC, etc.

Managed lifecycle

- bound to a node, restart in place
- can die, cannot be reborn with same ID

**Example: data puller & web server**



# Volumes

## Pod-scoped storage

Support many types of volume plugins

- Empty dir (and tmpfs)
- Host path
- Git repository
- GCE Persistent Disk
- AWS Elastic Block Store
- Azure File Storage
- iSCSI
- Flocker
- NFS
- vSphere
- GlusterFS
- Ceph File and RBD
- Cinder
- FibreChannel
- Secret, ConfigMap, DownwardAPI
- Flex (exec a binary)
- ...



# Labels & Selectors



# Labels

Arbitrary metadata

Attached to any API object


Generally represent **identity**

Queryable by **selectors**





- think SQL 'select ... where ...'

The **only** grouping mechanism

- pods under a ReplicaSet
- pods in a Service
- capabilities of a node (constraints)



# Selectors



App: MyApp

Phase: prod


Role: BE

App: MyApp

Phase: test


Role: BE

# Selectors




**App = MyApp**

# Selectors




**App = MyApp, Role = FE**

# Selectors




**App = MyApp, Role = BE**

# Selectors



**App = MyApp, Phase = prod**

# Selectors



**App = MyApp, Phase = test**

# Replication



# ReplicaSets

A simple control loop

Runs out-of-process wrt API server

**One job:** ensure N copies of a pod

- grouped by a selector
- too few? start some
- too many? kill some


Layered on top of the public Pod API

Replicated pods are **fungible**

- No implied order or identity

## ReplicaSet

- **name** = "my-rc"
- **selector** = {"App": "MyApp"}
- **template** = { ... }
- **replicas** = 4



API Server

# Control loops: the Reconciler Pattern

Drive **current state** -> **desired state**


Act independently

APIs - **no shortcuts** or back doors

Observed state is truth\*

Recurring pattern in the system

**Example: ReplicaSet**



\* Observations are really stale caches of what once was your view of truth.

# Services



# Services

A group of **pods** that work together

- grouped by a selector

Defines access policy

- “load balanced” or “headless”


Can have a stable **virtual IP** and port

- also a DNS name


VIP is managed by *kube-proxy*

- watches all services
- updates iptables when backends change
- default implementation - can be replaced!


Hides complexity




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy




# iptables kube-proxy



# iptables kube-proxy



# iptables kube-proxy



# External services

Services VIPs are only available **inside** the cluster

Need to receive traffic from “the outside world”


Service “type”

- NodePort: expose on a port on every node
- LoadBalancer: provision a cloud load-balancer

DiY load-balancer solutions

- socat (for nodePort remapping)
- haproxy
- nginx

Ingress (L7 LB)



# Ingress (L7 LB)

Many apps are HTTP/HTTPS

Services are L4 (IP + port)


Ingress maps incoming traffic to backend services

- by HTTP host headers
- by HTTP URL paths

HAProxy, NGINX, AWS and GCE implementations in progress

Now with SSL!

Status: BETA in Kubernetes v1.2




# Rolling Update



# Rolling Update

## ReplicaSet


- name: my-app-v1
- replicas: 3
- selector:
  - app: MyApp
  - version: v1



# Rolling Update

## ReplicaSet

- name: my-app-v1
- replicas: 3
- selector:
  - app: MyApp
  - version: v1




## ReplicaSet

- name: my-app-v2
- replicas: 0
- selector:
  - app: MyApp
  - version: v2

# Rolling Update

## ReplicaSet

- name: my-app-v1
- replicas: 3
- selector:
  - app: MyApp
  - version: v1




## ReplicaSet

- name: my-app-v2
- replicas: 1
- selector:
  - app: MyApp
  - version: v2

# Rolling Update

## ReplicaSet

- name: my-app-v1
- replicas: 2
- selector:
  - app: MyApp
  - version: v1




## ReplicaSet

- name: my-app-v2
- replicas: 1
- selector:
  - app: MyApp
  - version: v2

# Rolling Update

## ReplicaSet

- name: my-app-v1
- replicas: 2
- selector:
  - app: MyApp
  - version: v1



## ReplicaSet

- name: my-app-v2
- replicas: 2
- selector:
  - app: MyApp
  - version: v2


# Rolling Update

## ReplicaSet

- name: my-app-v1
- replicas: 1
- selector:
  - app: MyApp
  - version: v1




**Service**  
- app: MyApp



## ReplicaSet


- name: my-app-v2
- replicas: 2
- selector:
  - app: MyApp
  - version: v2



# Rolling Update


## ReplicaSet

- name: my-app-v1
- replicas: 1
- selector:
  - app: MyApp
  - version: v1




## Service

- app: MyApp



## ReplicaSet

- name: my-app-v2
- replicas: 3
- selector:
  - app: MyApp
  - version: v2



# Rolling Update


**ReplicaSet**  
- name: my-app-v1  
- replicas: 0  
- selector:  
  - app: MyApp  
  - version: v1




**Service**  
- app: MyApp



**ReplicaSet**  
- name: my-app-v2  
- replicas: 3  
- selector:  
  - app: MyApp  
  - version: v2



# Rolling Update



# Deployments



# Deployments

## Updates-as-a-service

- Rolling update is imperative, client-side


Deployment manages replica changes for you

- stable object name
- updates are configurable, done server-side
- `kubectl edit` or `kubectl apply`

Aggregates stats

Can have multiple updates in flight

**Status: BETA in Kubernetes v1.2**



# DaemonSets



# DaemonSets

**Problem: how to run a Pod on every node?**

- or a subset of nodes


Similar to ReplicaSet

- principle: do one thing, don't overload

“Which nodes?” is a selector

Use familiar tools and patterns

**Status: BETA in Kubernetes v1.2**



# Jobs



# Jobs


**Run-to-completion**, as opposed to run-forever

- Express parallelism vs. required completions
- Workflow: restart on failure
- Build/test: don't restart on failure

Aggregates success/failure counts

Built for batch and big-data work

**Status: GA in Kubernetes v1.2**



# PersistentVolumes



# PersistentVolumes

A higher-level storage abstraction

- insulation from any one cloud environment


Admin provisions them, users claim them

- **NEW: auto-provisioning (alpha in v1.2)**


Independent lifetime from consumers

- lives until user is done with it
- can be handed-off between pods

Dynamically “scheduled” and managed, like nodes and pods




# PersistentVolumes




**Cluster  
Admin**



# PersistentVolumes




# PersistentVolumes




**Cluster  
Admin**




**User**




# PersistentVolumes




# PersistentVolumes




# PersistentVolumes




# PersistentVolumes



# PersistentVolumes




# PersistentVolumes




User


# PersistentVolumes




# PersistentVolumes




# PersistentVolumes



# PersistentVolumes




# PersistentVolumes



**Cluster  
Admin**



**User**



**Recycler**

# StatefulSets



# StatefulSets


## Goal: enable clustered software on Kubernetes

- mysql, redis, zookeeper, ...

Clustered apps need “identity” and sequencing guarantees

- stable hostname, available in DNS
- an ordinal index
- stable storage: linked to the ordinal & hostname
- discovery of peers for quorum
- startup/teardown ordering

Status: ALPHA in Kubernetes v1.3



# ConfigMaps



# ConfigMaps

## Goal: manage app configuration

- ...without making overly-brittle container images

[12-factor](#) says config comes from the environment


- Kubernetes is the environment

Manage config via the Kubernetes API

Inject config as a virtual volume into your Pods

- late-binding, live-updated (atomic)
- also available as env vars

Status: GA in Kubernetes v1.2



# Secrets



# Secrets

**Goal: grant a pod access to a secured *something***

- don't put secrets in the container image!


[12-factor](#) says config comes from the environment

- Kubernetes is the environment

Manage secrets via the Kubernetes API

Inject secrets as virtual volumes into your Pods

- late-binding, tmpfs - never touches disk
- also available as env vars



# HorizontalPodAutoscalers



# HorizontalPodAutoScalers

**Goal: Automatically scale pods as needed**


- based on CPU utilization (for now)
- custom metrics in Alpha

Efficiency now, capacity when you need it

Operates within user-defined min/max bounds

Set it and forget it

**Status: GA in Kubernetes v1.2**



# Multi-Zone Clusters



# Multi-Zone Clusters

## Goal: zone-fault tolerance for applications


Zero API changes relative to kubernetes

- Create services, ReplicaSets, etc. exactly as usual

Nodes and PersistentVolumes are labelled with their availability zone

- Fully automatic for GKE, GCE, AWS
- Manual for on-premise and other cloud providers (for now)

**Status: GA in Kubernetes v1.2**



# Namespaces



# Namespaces

## Problem: I have too much stuff!

- name collisions in the API
- poor isolation between users
- don't want to expose things like Secrets

## Solution: Slice up the cluster

- create new Namespaces as needed
  - per-user, per-app, per-department, etc.
- part of the API - NOT private machines
- most API objects are namespaced
  - part of the REST URL path
- Namespaces are just another API object
- One-step cleanup - delete the Namespace
- Obvious hook for policy enforcement (e.g. quota)




# Resource Isolation



# Resource Isolation

## Principles:

- Apps must not be able to affect each other's performance
  - if so it is an **isolation failure**
- Repeated runs of the same app should see ~equal behavior
- QoS levels drives resource decisions in (soft) real-time
- Correct in all cases, optimal in some
  - reduce unreliable components
- SLOs are the lingua franca



# Strong isolation

## Pros:

- Sharing - users don't worry about interference (aka the noisy neighbor problem)
- Predictable - allows us to offer strong SLAs to apps

## Cons:

- Stranding - arbitrary slices mean some resources get lost
- Confusing - how do I know how much I need?
  - analog: what size VM should I use?
  - smart auto-scaling is needed!
- Expensive - you pay for certainty

In reality this is a multi-dimensional bin-packing problem: CPU, memory, disk space, IO bandwidth, network bandwidth, ...

# Requests and Limits

## Request:

- how much of a resource you are asking to use, with a strong guarantee of availability
  - CPU (seconds/second)
  - RAM (bytes)
- scheduler will not over-commit requests

## Limit:

- max amount of a resource you can access

## Repercussions:

- Usage > Request: resources **might** be available
- Usage > Limit: throttled or killed



# Quality of Service

Defined in terms of Request and Limit

**Guaranteed:** highest protection

- $\text{request} > 0 \ \&\& \ \text{limit} == \text{request}$

**Burstable:** medium protection


- $\text{request} > 0 \ \&\& \ \text{limit} > \text{request}$

**Best Effort:** lowest protection

- $\text{request} == 0$

What does “protection” mean?

- OOM score
- CPU scheduling



# Quota and Limits



# ResourceQuota

Admission control: apply limits in **aggregate**

**Per-namespace:** ensure no user/app/department abuses the cluster

Reminiscent of disk quota by design

Applies to each type of resource

- CPU and memory for now

Disallows pods without resources



# LimitRange

Admission control: limit the limits

- min and max
- ratio of limit/request

**Default values for unspecified limits**

**Per-namespace**

Together with ResourceQuota gives cluster admins powerful tools



# Cluster Auto-Scaling



# Cluster Autoscaler

Add nodes when needed

- there are pending pods
- some pending pods would fit if we add a node

Remove nodes when not needed

- after removal, all pods must fit remaining nodes

Status: Works on GCE, GKE and AWS



# Scalability




# Scalability & Performance

SLO met at <2000 nodes, <60000 pods

- 99% of API calls return in < 1 second
- 99% of pods start in < 5 seconds

Coming soon

- protobufs in API storage (already enabled on the wire)
- 5000 nodes



# Design principles

**Declarative > imperative:** State your desired results, let the system actuate

**Control loops:** Observe, rectify, repeat

**Simple > Complex:** Try to do as little as possible

**Modularity:** Components, interfaces, & plugins

**Legacy compatible:** Requiring apps to change is a non-starter

**Network-centric:** IP addresses are cheap

**No grouping:** Labels are the only groups

**Sets > Pets:** Manage your workload in bulk

**Open > Closed:** Open Source, standards, REST, JSON, etc.

# Kubernetes (K8s) Community

~5k Commits  
in 1.4 over 3  
months

> 800 Unique  
Contributors

Top 0.01% of  
all Github  
Projects


2500+ External  
Projects Based  
on K8s

## Companies Contributing



## Companies Using





"Niantic chose GKE for its ability to **orchestrate their container cluster at planetary-scale**, freeing its team to focus on deploying live changes for their players." - Niantic

# Further Reading

If this talk was interesting, deeper academic reading on cluster management:

## **“Borg, Omega, and Kubernetes”**

ACM Queue, March 2, 2016, [Volume 14, issue 1](#)

<http://queue.acm.org/detail.cfm?id=2898444>

Or a hands-on “Hello World” quickstart to build a Docker image and run it on a Kubernetes cluster:

<http://kubernetes.io/docs/hellonode/>

Another hard problem: how do you run  $N$  Kubernetes clusters as a service?

- create/delete, update, monitor, repair, escalate, upgrade, backup/restore, zonal isolation, incremental rollouts, support ticket escalation, provisioning, and more!



# Questions?

Potential discussion:

- What about Docker Swarm?
- ... Mesos?
- What's next for Kubernetes and Container Engine?
- Why Google not FB/Uber/MS/Ama/etc?
- How do I get an internship / job?
  - Let's discuss!

- Alex on Philosophy:
  - Imperative vs. declarative
  - Orchestration vs. choreography
  - Product vs. tech
  - User guide vs. design doc
  - Engineering code vs. organizations
  - Your team is a design parameter
  - Launch and iterate; MVP

More questions?

Happy to chat!

- Lunch
- 1:1's after that
- [mohr@google.com](mailto:mohr@google.com)
- [590s@alexmohr.com](mailto:590s@alexmohr.com)